

Columbia University

Driving Scalable Growth: Database Solutions for ABC Foodmart

Adeel Arif, Depali Kulkarni, Suchakrey Nitisanon,
Kristen Campbell, Sana Khan, Malaikah Khan



The Client & Challenge

ABC Foodmart

Today

A grocery chain operating two locations in Queens, NY, preparing to expand into Brooklyn with three additional stores.

Current State: Relying on large spreadsheets for inventory, staffing, and vendor management across locations.

The Problem

Operational Inefficiencies:

- Error-prone bulk updates
- Weak auditability and tracking
- No cross-location coordination
- Limited analytical capabilities

Multi-borough expansion requires centralized data infrastructure

Our Solution: A Centralized Relational Database



Six Interconnected Tables

Stores, customers, products, aisles, transactions, and transaction items—all normalized to 3NF standards.



Operational Efficiency

Replace spreadsheet errors, streamline inventory management, and coordinate across all locations seamlessly.



Customer Loyalty System

Track purchasing behavior, reward points, and transaction history to build lasting customer relationships.

This infrastructure supports ABC Foodmart's expansion while creating a foundation for data-driven decision making and customer engagement.

Sample Data

grocery_chain_data.csv (163.33 kB)

Detail Compact Column

10 of 11 columns ▾

customer_id	store_name	transaction_date	aisle	product_name	# quantity	# unit_price
2824	GreenGrocer Plaza	2023-08-26	Produce	Pasta	2	7.46
5506	ValuePlus Market	2024-02-13	Dairy	Cheese	1	1.85
4657	ValuePlus Market	2023-11-23	Bakery	Onions	4	7.38
2679	SuperSave Central	2025-01-13	Snacks & Candy	Cereal	3	5.5
9935	GreenGrocer Plaza	2023-10-13	Canned Goods	Orange Juice	5	8.66
7912	Corner Grocery	2023-08-31	Dairy	Onions	3	25.04
1488	GreenGrocer Plaza	2023-10-18	Produce	Ground Beef	3	29.56
4582	City Fresh Store	2024-02-03	Health & Wellness	Orange Juice	2	11.73
9279	GreenGrocer Plaza	2024-11-16	Household Items	Apples	5	22.13
1434	QuickStop Market	2024-10-15	Health & Wellness	Onions	2	20.84
4257	FreshMart Downtown	2025-02-13	Canned Goods	Tomatoes	3	2.91

The dataset contains **1,980 rows** across **11 columns**, sourced from Kaggle's grocery store sales records. Key fields include `customer_id`, `store_name`, product details, quantity, `transaction_date`, and pricing information.

This data naturally maps to normalized tables supporting ABC Foodmart's expansion needs while enabling customer behavior analysis and business metrics tracking.

Data Source

[Kaggle Grocery Store](#)

[Sales Dataset](#)

Structured to support multi-store operations and analytical queries

Database Schema

Design

Design Principles

Our schema follows **Third Normal Form (3NF)** to minimize redundancy and ensure data integrity at scale.

01

Entity

Identification

customers, stores, products, aisles, transactions, transaction

items

02

Relationship

Mapping

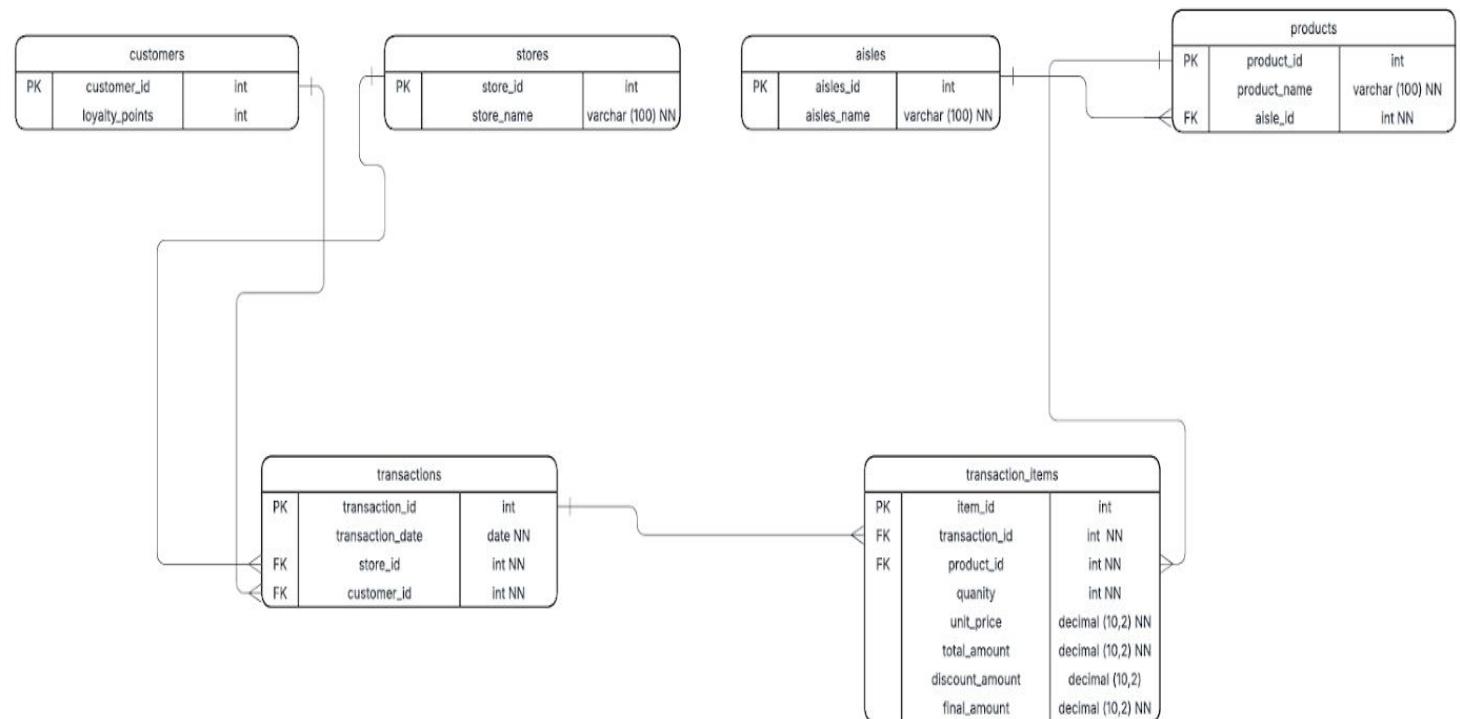
One-to-many relationships between core entities via foreign keys

03

Extensibility

Room to add staffing, vendor, and supply chain data without breaking existing structure

Entity Relationship Diagram showing table connections and foreign key relationships



The ETL Pipeline

1

Extract

Load raw Kaggle dataset into Pandas DataFrame for quality assessment

2

Transform

Correct product-to-aisle mappings, normalize data into six tables, preserve original customer IDs

3

Load

Use SQLAlchemy to insert data into PostgreSQL following foreign key dependency order



Critical Fix: Manual correction of incorrect product-aisle assignments (e.g., pasta was incorrectly categorized under produce) ensured data integrity before database loading.

Customer Needs: What &

Why For Analysts

- **Sales Performance**

Analysis

Understand which stores, products, and aisles drive revenue

- **Customer Behavior**

Insights

Track purchasing patterns, repeat visits, and loyalty metrics

- **Discount**

Effectiveness

Measure promotion usage on sales volume

- **Time-Based**

Trends

Identify daily, weekly, and seasonal sales patterns for inventory planning

For Executives

- **Real-Time Performance**

Monitoring

Track KPIs across all locations without technical expertise

- **Strategic Decision**

Support

Data-driven insights for staffing, inventory, and expansion planning

- **Operational**

Efficiency

Identify low performing areas and optimize resource allocation

- **Accessible**

Analytics

Visual dashboards replace complex SQL queries

Analytical Tools: The

How

1

PostgreSQL + pgAdmin4

For Analysts: Direct database access with 16 pre-built SQL scripts answering key business questions

- Store revenue analysis
- Product revenue rankings
- Customer spending patterns
- Discount impact assessment

2

Metabase Dashboards

Interactive

visualizations with no SQL knowledge required

- Real-time sales monitoring
- Aisle traffic visualization
- Store comparison metrics
- Automated email reports

3

Python ETL Scripts

Engineers: Automated data transformation and loading using Pandas and SQLAlchemy

- Data validation and cleaning
- Dependency-aware loading
- Error handling and logging

Analytical Procedures: Business Questions Answered Through SQL



Store Performance & Revenue Optimization

Understand how each store performs financially and operationally using key metrics such as:

Total revenue and revenue after discounts

Average order value (AOV) and transaction volume

Store-level discount usage and its impact on earnings



Product & Aisle Sales Intelligence

Understand what drives sales performance across product lines and categories using insights such as:

Top-revenue products and top-selling SKUs

Aisle-level revenue and units sold

Daily sales trends for inventory and demand planning



Customer Behavior & Loyalty Insights

Understand customer value and engagement patterns using insights such as:

Highest-spending and highest-value customers

Repeat visit frequency across locations

Lifetime spend patterns for loyalty and retention

SQL Insight Deep Dive: Post-Discount Revenue Analysis

Query Purpose

Calculate each store's total revenue after subtracting all discounts.

Key Fields Used

Store-Level Fields

`store_id, store_name`

Transaction Fields

`final_amount, discount_amount`

Logic Summary

Revenue After Discount =
`SUM(final_amount - discount_amount)`

store_name	revenue_after_discount
ValuePlus Market	15896.97
FreshMart Downtown	15131.58
GreenGrocer Plaza	14773.07
Corner Grocery	14198.28
MegaMart Westside	13187.62

What the Business Can Learn

Identify the **true** top-performing stores after removing discount inflation.

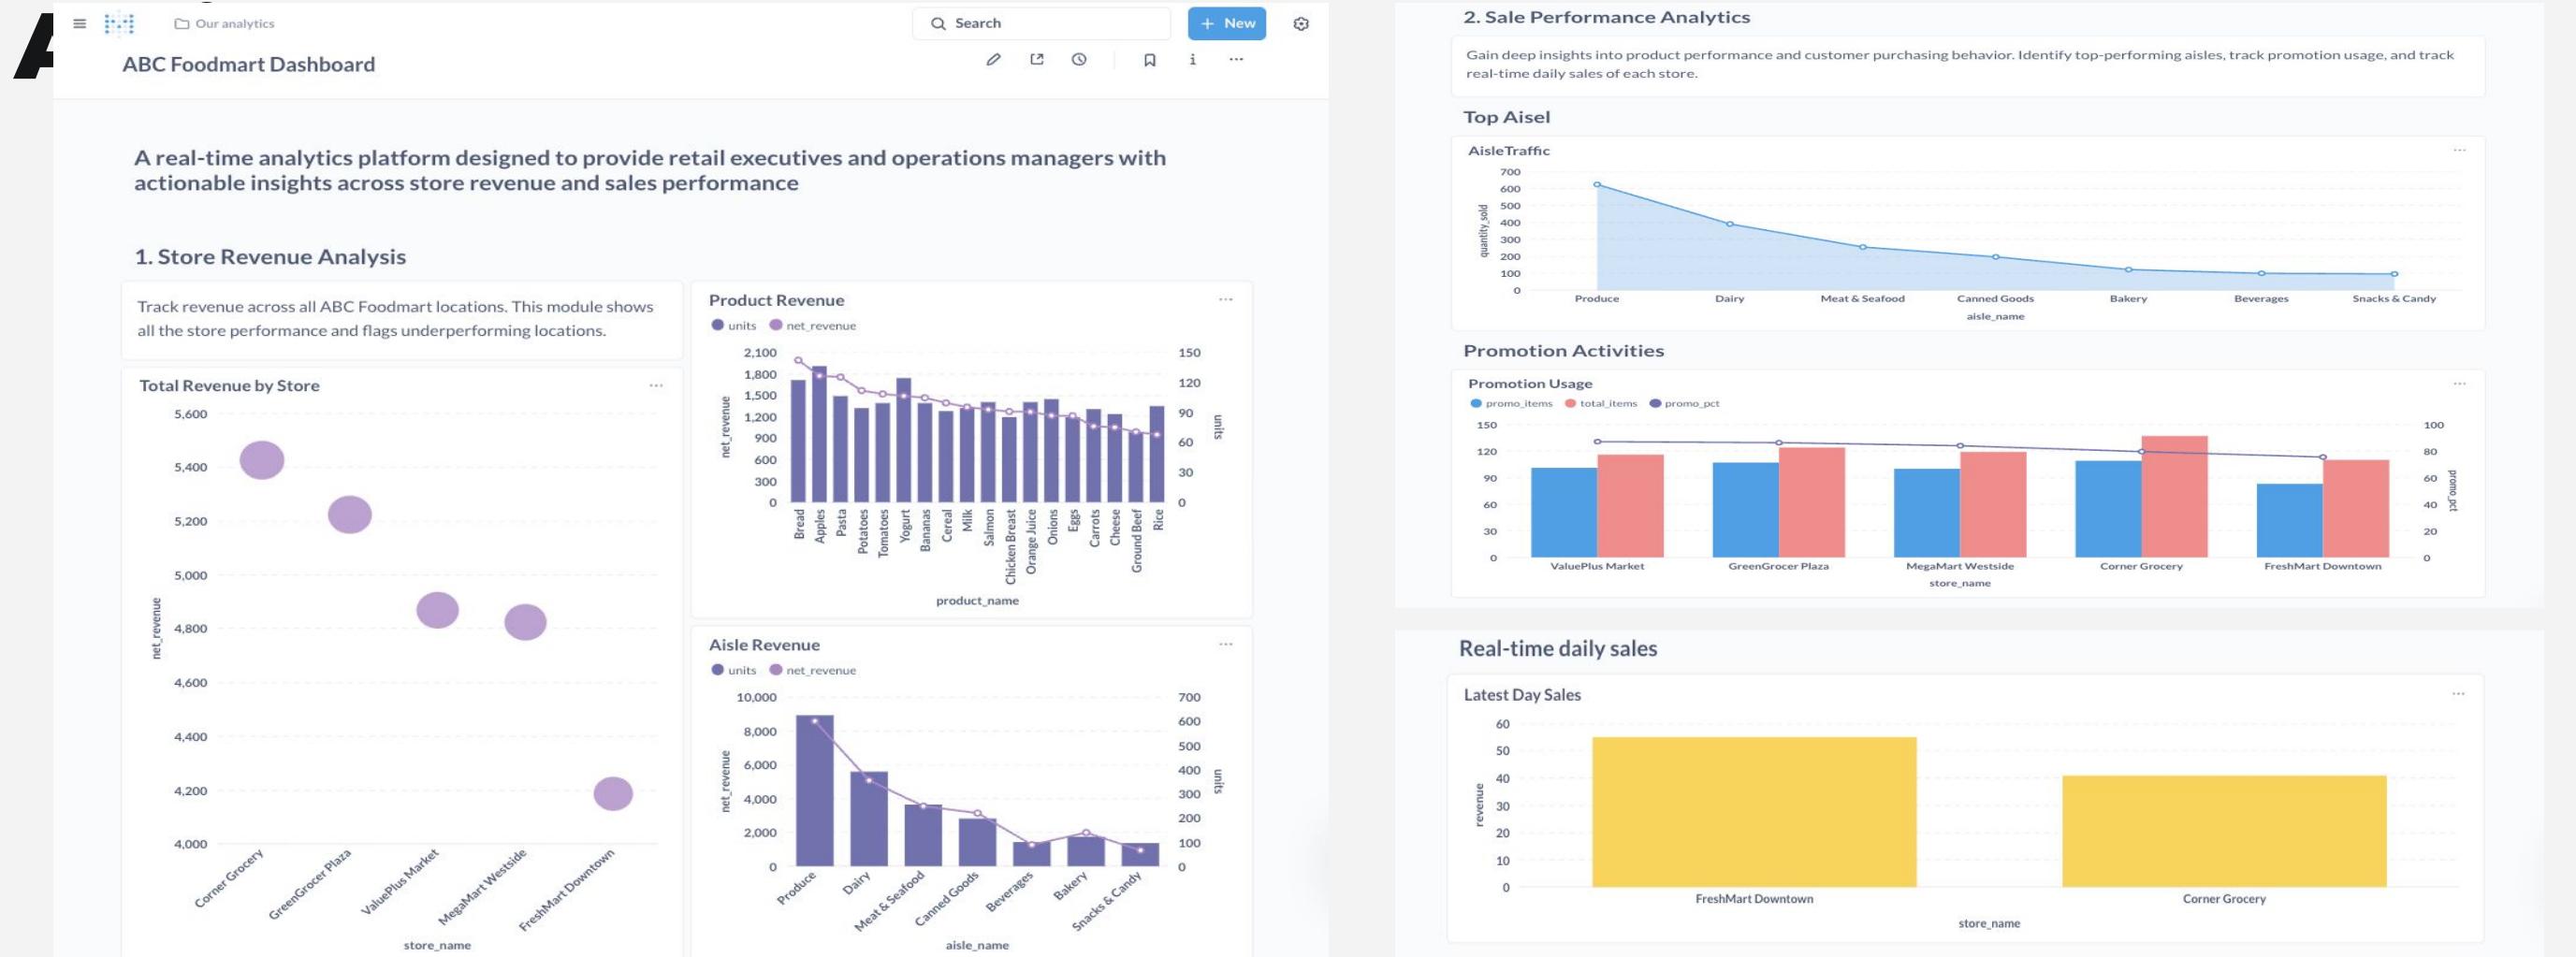
Understand whether current discount strategies are boosting sales or silently eroding margins.

Impact on Decision-Making

Leadership can redesign or reduce discount programs that aren't generating real revenue uplift.

High-performing stores (*post adjustment*) become models for operational best practices and resource allocation.

Dashboard Insights in



Daily Revenue

Monitor revenue performance in real-time with customizable date range, location, product, and aisle breakdowns

Aisle Traffic

Analyze traffic through high and low-traffic areas to optimize product placement and promotional strategies

Promotion

Effectiveness

Measure promotion usage by store with customized filters such as product and aisle to refine pricing strategies

Dashboard Demo

Impact & Next Steps

What We Delivered

- Replaced error-prone spreadsheets with a centralized PostgreSQL database following 3NF standards
- Built a Python ETL pipeline with data validation, transformation, and dependency-aware loading
- Created 16 SQL analytical procedures for deep-dive analysis of sales, customers, and products
- Designed executive dashboards in Metabase for real-time monitoring without technical skills

Business Impact

- **Scalability:** Infrastructure ready to support Brooklyn expansion and beyond
- **Data Integrity:** Elimination of spreadsheet errors and improved auditability
- **Faster Insights:** Real-time analytics enable agile decision-making
- **Customer Loyalty:** Transaction tracking powers rewards programs